

## Data Mining Assignment 4

1) Read Chapter 4 (all sections) and Chapter 5 (Sections 5.2, 5.5, 5.6 and 5.7).

2) Repeat In Class Exercise #38 using the misclassification error rate instead of information gain to determine the best split. Which of these splits considered is the best according to misclassification error rate?

A	B	Class Label
T	F	+
T	T	+
T	T	+
T	F	-
T	T	+
F	F	-
F	F	-
F	F	-
T	T	-
T	F	-

- As we can see from the above table, if we split on A, the misclassification error would be:  $3 / 10 = 0.3$ . Because in rows 4, 9 and 10 we can see that three records of A are misclassified and 10 is the total number of records.
- If we split on B, there are misclassifications in row 1 and 9 with respect to B, so the rate would be 0.2.
- Since the misclassification rate is low when we split the data set on B, we need to induct our decision tree based on B split.

3) Repeat In Class Exercise #39 using the misclassification error rate instead of information gain to determine the best split. Which of these splits considered is the best according to misclassification error rate?

Instance	$a_1$	$a_2$	$a_3$	Target Class
1	T	T	1.0	+
2	T	T	6.0	+
3	T	F	5.0	-
4	F	F	4.0	+
5	F	T	7.0	-
6	F	T	3.0	-
7	F	F	8.0	-
8	T	F	7.0	+
9	F	T	5.0	-

- Splitting on  $a_1$ , the misclassification error rate =  $2 / 9 = 0.22$
- Splitting on  $a_2$ , the misclassification error rate =  $5 / 9 = 0.55$
- Splitting on  $a_3$ , [ So, splitting on  $a_3$  will not be straight because it is not a nominal value or categorical value. Here, the  $a_3$  has discrete values and I decided to split on condition  $a_3 < 5.0$  as +  $a_3 \geq 5.0$  as -, the misclassification error rate would be =  $3 / 9 = 0.33$

4) The file [http://www-stat.wharton.upenn.edu/~dmease/rpart\\_text\\_example.txt](http://www-stat.wharton.upenn.edu/~dmease/rpart_text_example.txt) gives an example of text output for a tree fit using the `rpart()` function in R from the library `rpart`. Use this tree to predict the class labels for the 10 observations in the test data [http://www-stat.wharton.upenn.edu/~dmease/test\\_data.csv](http://www-stat.wharton.upenn.edu/~dmease/test_data.csv) linked here. Do this manually - do not use R or any software.

- Age = Middle, Number = 5 and Start = 10, the class label is present, as we traverse from 1 -> 2 -> 5 -> 11
- Age = young, Number = 2, Start = 17, the class label is absent, as we traverse from 1 -> 2 -> 4 -> 8
- Age = old, Number = 10, Start = 6, the class label is present, as we traverse from 1 -> 3 -> 7 -> 15
- Age = young, Number = 2, Start = 17, the class label is absent, as we traverse from 1 -> 2 -> 4 -> 8

- Age = old, Number = 4, Start = 15, the class label is absent, as we traverse from 1 -> 2 -> 4 -> 8
- Age = middle, Number = 5, Start = 15, the class label is absent, as we traverse from 1 -> 2 -> 5 -> 10
- Age = young, Number = 3, Start = 13, the class label is absent, as we traverse from 1 -> 2 -> 4 -> 9
- Age = old, Number = 5, Start = 8, the class label is present, as we traverse from 1 -> 3 -> 7 -> 15
- Age = young, Number = 7, Start = 9, the class label is absent, as we traverse from 1 -> 2 -> 4 -> 9
- Age = middle, Number = 3, Start = 13, the class label is absent, as we traverse from 1 -> 2 -> 5 -> 10

5) I split the popular sonar data set into a training set ([http://www-stat.wharton.upenn.edu/~dmease/sonar\\_train.csv](http://www-stat.wharton.upenn.edu/~dmease/sonar_train.csv)) and a test set ([http://www-stat.wharton.upenn.edu/~dmease/sonar\\_test.csv](http://www-stat.wharton.upenn.edu/~dmease/sonar_test.csv)). Use R to compute the misclassification error rate on the test set when training on the training set for a tree of depth 5 using all the default values except `control=rpart.control(minsplit=0,minbucket=0,cp=-1, maxcompete=0, maxsurrogate=0, usesurrogate=0, xval=0,maxdepth=5)`. Remember that the 61st column is the response and the other 60 columns are the predictor

```
Console Terminal x Jobs x
F:/Data Science/DataScience_2019501111/Data Mining/DM Assignment4/
> setwd("F:\\Data Science\\DataScience_2019501111\\Data Mining\\DM Assignment4")
> getwd()
[1] "F:/Data Science/DataScience_2019501111/Data Mining/DM Assignment4"
>
> train = read.csv("sonar_train.csv",header = FALSE)
> test = read.csv("sonar_test.csv",header = FALSE)
> summary(train)
```

v1		v2		v3		v4	
Min.	:0.00150	Min.	:0.00170	Min.	:0.00150	Min.	:0.00580
1st Qu.	:0.01523	1st Qu.	:0.01675	1st Qu.	:0.01903	1st Qu.	:0.02390
Median	:0.02285	Median	:0.03040	Median	:0.03415	Median	:0.03905
Mean	:0.02868	Mean	:0.03677	Mean	:0.04250	Mean	:0.05127
3rd Qu.	:0.03445	3rd Qu.	:0.04735	3rd Qu.	:0.05375	3rd Qu.	:0.06270
Max.	:0.13710	Max.	:0.16320	Max.	:0.19970	Max.	:0.26040

v5		v6		v7		v8	
Min.	:0.00670	Min.	:0.01160	Min.	:0.0033	Min.	:0.00980
1st Qu.	:0.03955	1st Qu.	:0.06310	1st Qu.	:0.0808	1st Qu.	:0.07485
Median	:0.06090	Median	:0.08695	Median	:0.1055	Median	:0.11405
Mean	:0.07355	Mean	:0.10467	Mean	:0.1222	Mean	:0.13749
3rd Qu.	:0.09625	3rd Qu.	:0.13378	3rd Qu.	:0.1568	3rd Qu.	:0.17095
Max.	:0.32250	Max.	:0.38230	Max.	:0.3729	Max.	:0.45900

v9		v10		v11		v12	
Min.	:0.0155	Min.	:0.0242	Min.	:0.0327	Min.	:0.0269
1st Qu.	:0.0953	1st Qu.	:0.1057	1st Qu.	:0.1217	1st Qu.	:0.1211
Median	:0.1463	Median	:0.1736	Median	:0.2204	Median	:0.2493
Mean	:0.1839	Mean	:0.2123	Mean	:0.2399	Mean	:0.2542
3rd Qu.	:0.2455	3rd Qu.	:0.2774	3rd Qu.	:0.3128	3rd Qu.	:0.3444
Max.	:0.6828	Max.	:0.7106	Max.	:0.7342	Max.	:0.7060

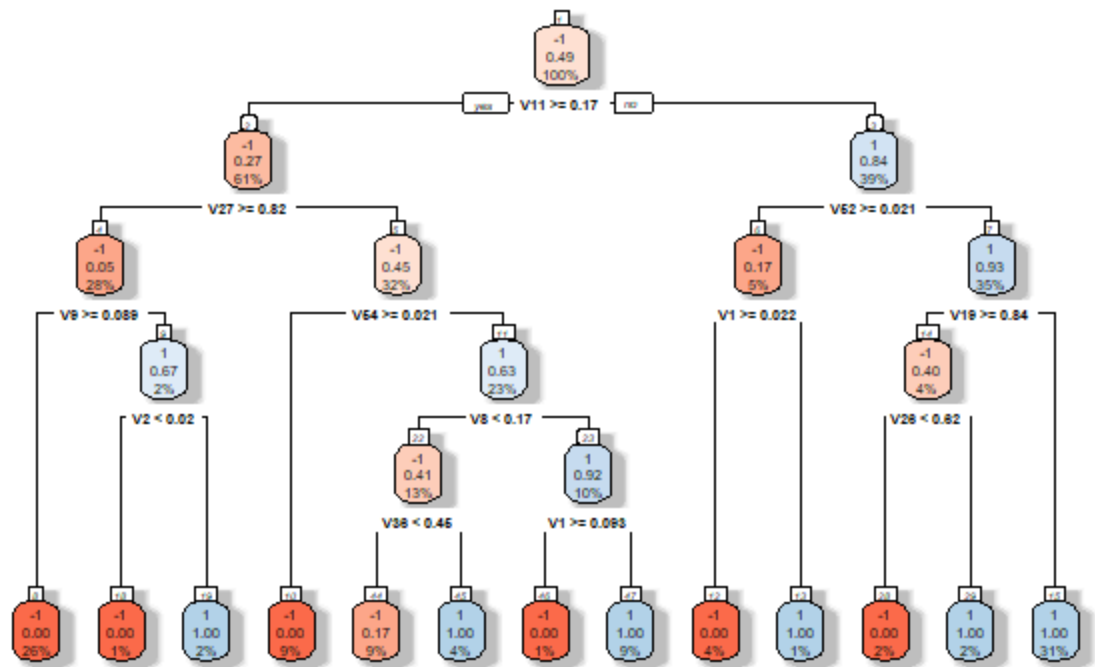
v13		v14		v15		v16	
Min.	:0.0252	Min.	:0.0336	Min.	:0.0031	Min.	:0.0162
1st Qu.	:0.1548	1st Qu.	:0.1568	1st Qu.	:0.1459	1st Qu.	:0.1805
Median	:0.2451	Median	:0.2614	Median	:0.2594	Median	:0.2903
Mean	:0.2723	Mean	:0.2809	Mean	:0.2959	Mean	:0.3549
3rd Qu.	:0.3679	3rd Qu.	:0.3877	3rd Qu.	:0.4291	3rd Qu.	:0.5059
Max.	:0.7131	Max.	:0.7842	Max.	:0.8392	Max.	:0.9444

v17		v18		v19		v20	
Min.	:0.0349	Min.	:0.0375	Min.	:0.0494	Min.	:0.0740
1st Qu.	:0.1940	1st Qu.	:0.2261	1st Qu.	:0.3119	1st Qu.	:0.3430
Median	:0.3036	Median	:0.3654	Median	:0.4433	Median	:0.5736

```

> dim(train)
[1] 130 61
> summary(test)
      v1      v2      v3      v4
Min.   :0.00250 Min.   :0.00060 Min.   :0.00510 Min.   :0.00610
1st Qu.:0.01268 1st Qu.:0.01560 1st Qu.:0.01843 1st Qu.:0.02898
Median :0.02245 Median :0.03085 Median :0.03450 Median :0.04830
Mean   :0.02996 Mean   :0.04121 Mean   :0.04606 Mean   :0.05826
3rd Qu.:0.03717 3rd Qu.:0.04893 3rd Qu.:0.06137 3rd Qu.:0.07093
Max.   :0.13130 Max.   :0.23390 Max.   :0.30590 Max.   :0.42640
      v5      v6      v7      v8
Min.   :0.00760 Min.   :0.01020 Min.   :0.01820 Min.   :0.0055
1st Qu.:0.03535 1st Qu.:0.07445 1st Qu.:0.08192 1st Qu.:0.0888
Median :0.06615 Median :0.09435 Median :0.11090 Median :0.1086
Mean   :0.07796 Mean   :0.10440 Mean   :0.12103 Mean   :0.1303
3rd Qu.:0.10743 3rd Qu.:0.13350 3rd Qu.:0.14822 3rd Qu.:0.1659
Max.   :0.40100 Max.   :0.25870 Max.   :0.30160 Max.   :0.4223
      v9      v10     v11     v12
Min.   :0.0075 Min.   :0.0113 Min.   :0.0289 Min.   :0.0236
1st Qu.:0.1032 1st Qu.:0.1251 1st Qu.:0.1475 1st Qu.:0.1565
Median :0.1560 Median :0.1888 Median :0.2354 Median :0.2490
Mean   :0.1682 Mean   :0.2016 Mean   :0.2295 Mean   :0.2436
3rd Qu.:0.2099 3rd Qu.:0.2541 3rd Qu.:0.2935 3rd Qu.:0.3105
Max.   :0.5744 Max.   :0.5378 Max.   :0.5533 Max.   :0.5771
      v13     v14     v15     v16
Min.   :0.0184 Min.   :0.0273 Min.   :0.0456 Min.   :0.0906
1st Qu.:0.1798 1st Qu.:0.2273 1st Qu.:0.1882 1st Qu.:0.2062
Median :0.2671 Median :0.2995 Median :0.3230 Median :0.3533
Max.   :0.53200 Max.   :0.043900 Max.   :1.0000
> dim(test)
[1] 78 61
>
> library(rpart)
> library(rpart.plot)
> help("rpart.control")
> help("rpart.plot")
> x <- train[,1:60]
> y <- as.factor(train[,61])
> model <- rpart(y~.,x,control=rpart.control(minsplit=0,minbucket=0,cp=-1,
axcompete=0,maxsurrogate=0,usesurrogate=0,xval=0,maxdepth=5))
> rpart.plot(model, box.palette="RdBu", shadow.col="gray", nn=TRUE)
>
> x_test <- test[,1:60]
> y_test <- as.factor(test[,61])
> 1 - sum(y_test == predict(model,x_test,type = "class")) / length(y_test)
[1] 0.2564103
>

```



6) Do Chapter 5 textbook problem #17 (parts a and c only) on pages 322-323. Note that there is a typo in part c - it should read "Repeat the analysis for part (b)". We will do part b in class.

You are asked to evaluate the performance of two classification models, M1 and M2. The test set you have chosen contains 26 binary attributes, labeled as A through Z.

Table 5.14 shows the posterior probabilities obtained by applying the models to the test set. (Only the posterior probabilities for the positive class are shown). As this is a two-class problem,  $P(-) = 1 - P(+)$  and  $P(-|A, \dots, Z) = 1 - P(+|A, \dots, Z)$ . Assume that we are mostly interested in detecting instances from the positive class.

**Table 5.14.** Posterior probabilities for Exercise 17.

Instance	True Class	$P(+ A, \dots, Z, M_1)$	$P(+ A, \dots, Z, M_2)$
1	+	0.73	0.61
2	+	0.69	0.03
3	-	0.44	0.68
4	-	0.55	0.31
5	+	0.67	0.45
6	+	0.47	0.09
7	-	0.08	0.38
8	-	0.15	0.05
9	+	0.45	0.01
10	-	0.35	0.04

(a) Plot the ROC curve for both M1 and M2. (You should plot them on the same graph.) Which model do you think is better? Explain your reasons.

A) From the above figure we can see that the M1 model is better as the TPR is more than that of the M2.

(c) Plot the ROC curve for both M1 and M2. (You should plot them on the same graph.) Which model do you think is better? Explain your reasons.

A) For model M2: Precision =  $1/2 = 50\%$ . Recall =  $1/5 = 20\%$ . F-measure =  $(2 \times .5 \times .2) / (.5 + .2) = 0.2857$ .

7) Compute the misclassification error on the training data for the Random Forest classifier from In Class Exercise #47. Show your R code for doing this.

```
Console Terminal x Jobs x
F:/Data Science/DataScience_2019501111/Data Mining/DM Assignment4/
> install.packages("randomForest")
WARNING: Rtools is required to build R packages but is not currently installed. Please do
wnload and install the appropriate version of Rtools before proceeding:

https://cran.rstudio.com/bin/windows/Rtools/
Installing package into 'C:/Users/hp/Documents/R/win-library/4.0'
(as 'lib' is unspecified)
trying URL 'https://cran.rstudio.com/bin/windows/contrib/4.0/randomForest_4.6-14.zip'
Content type 'application/zip' length 249367 bytes (243 KB)
downloaded 243 KB

package 'randomForest' successfully unpacked and MD5 sums checked

The downloaded binary packages are in
C:\Users\hp\AppData\Local\Temp\RtmpiHeOT\downloaded_packages
> library("randomForest")
randomForest 4.6-14
Type rfNews() to see new features/changes/bug fixes.
Warning message:
package 'randomForest' was built under R version 4.0.3
>
> train <- read.csv("sonar_test.csv", header = FALSE)
> test <- read.csv("sonar_test.csv", header = FALSE)
>
> x_train = train[,1:60]
> y_train = as.factor(train[,61])
>
> x_test = test[,1:60]
> y_test = as.factor(test[,61])
>
> model<-randomForest(x_train, y_train)
> 1 - sum(y_train == predict(model, x_train)) / length(y_train)
[1] 0
> |
```

8) This question deals with In Class Exercise #42.

a) Repeat In Class Exercise #42 for the k-nearest neighbor classifier for  $k=5$  and  $k=6$ .



```
Console Terminal x Jobs x
F:/Data Science/DataScience_2019501111/Data Mining/DM Assignment4/
> library(class)
>
> train <- read.csv("sonar_test.csv", header = FALSE)
> test <- read.csv("sonar_test.csv", header = FALSE)
>
> x_train = train[,1:60]
> y_train = as.factor(train[,61])
>
> x_test = test[,1:60]
> y_test = as.factor(test[,61])
>
> help("knn")
> model1<-knn(x_train, x_test,y_train, k = 5)
> 1 - sum(y_test == model1) / length(y_test)
[1] 0.2051282
>
> model2<-knn(x_train, x_test, y_train, k = 6)
> 1 - sum(y_test == model2) / length(y_test)
[1] 0.2692308
```

b) Repeat part a using the exact same R code a few times. Explain why both the training errors and the test errors often change for k=6 but not for k=5. Hint: Read the help on the knn function if you do not know.

```
> setwd("F:\\Data Science\\DataScience_2019501111\\Data Mining\\DM Assignment4")
> library(class)
>
> train <- read.csv("sonar_test.csv", header = FALSE)
> test <- read.csv("sonar_test.csv", header = FALSE)
>
> x_train = train[,1:60]
> y_train = as.factor(train[,61])
>
> x_test = test[,1:60]
> y_test = as.factor(test[,61])
>
> help("knn")
> model1<-knn(x_train, x_test,y_train, k = 5)
> 1 - sum(y_test == model1) / length(y_test)
[1] 0.2051282
>
> model2<-knn(x_train, x_test, y_train, k = 6)
> 1 - sum(y_test == model2) / length(y_test)
[1] 0.3461538
> |
```