

Data Mining- Lab Exam

Time: 24 hours

Marks:100

Open a document and update document with your answers for each question and submit it.

1. a) For the dataset BSE_Sensex_Index.csv, create an extra column of successive differences for each column of numeric values in this data file. Extract two simple random samples with replacement of 1000 and 3000 observations (rows). Show your R commands for doing this. Do the same thing by using Excel. Show your Excel commands.

Note: Successive difference for date d1= (date d1 value-immediate available previous date of d1 value)/immediate available previous date of d1. For the last row fill up values with mean of its immediate three previous row values.

Console			
Terminal x Jobs x			
F:\Data Science\DataScience_2019501111\Data Mining\Final exam/			
<pre> open High Low Close Min. : 17.08 Min. : 17.08 Min. : 17.08 Min. : 17.08 1st Qu.: 83.43 1st Qu.: 84.07 1st Qu.: 82.50 1st Qu.: 83.17 Median : 116.45 Median : 117.59 Median : 115.03 Median : 116.34 Mean : 398.28 Mean : 401.03 Mean : 395.52 Mean : 398.50 3rd Qu.: 650.67 3rd Qu.: 654.12 3rd Qu.: 644.93 3rd Qu.: 648.62 Max. :1522.19 Max. :1526.45 Max. :1500.74 Max. :1506.34 volume Adj.Close open_new high_new Min. :7.800e+05 Min. : 17.08 Min. : -0.0582780 Min. : -0.0432817 1st Qu.:9.030e+06 1st Qu.: 83.17 1st Qu.: -0.0039618 1st Qu.: -0.0034432 Median :4.390e+07 Median : 116.34 Median : 0.0005554 Median : 0.0003948 Mean :5.964e+08 Mean : 398.50 Mean : 0.0005955 Mean : 0.0004185 3rd Qu.:4.035e+08 3rd Qu.: 648.62 3rd Qu.: 0.0050955 3rd Qu.: 0.0045302 Max. :8.926e+09 Max. :1506.34 Max. : 0.1067121 Max. : 0.0343908 low_new close_new volume_new Adj.close_new Min. : -0.0474458 Min. : -0.0402908 Min. : -0.718888 Min. : -0.0402908 1st Qu.: -0.0038973 1st Qu.: -0.0042513 1st Qu.: -0.105633 1st Qu.: -0.0042513 Median : 0.0008122 Median : 0.0003301 Median : -0.002597 Median : 0.0003301 Mean : 0.0005022 Mean : 0.0003370 Mean : 0.007552 Mean : 0.0003370 3rd Qu.: 0.0047861 3rd Qu.: 0.0048696 3rd Qu.: 0.103772 3rd Qu.: 0.0048696 Max. : 0.0910833 Max. : 0.0573273 Max. : 1.677175 Max. : 0.0573273 > data_3000 = randomRows(data, 3000) > > summary(data_3000) open High Low Close Min. : 16.72 Min. : 16.72 Min. : 16.72 Min. : 16.72 1st Qu.: 79.61 1st Qu.: 80.10 1st Qu.: 78.94 1st Qu.: 79.42 Median : 113.11 Median : 114.21 Median : 111.98 Median : 112.88 Mean : 379.96 Mean : 382.57 Mean : 377.36 Mean : 380.19 3rd Qu.: 495.77 3rd Qu.: 497.82 3rd Qu.: 494.57 3rd Qu.: 497.14 Max. :1556.51 Max. :1563.03 Max. :1554.09 Max. :1561.80 volume Adj.Close open_new high_new Min. :7.400e+05 Min. : 16.72 Min. : -0.0871188 Min. : -0.0685302 1st Qu.:5.972e+06 1st Qu.: 79.42 1st Qu.: -0.0039658 1st Qu.: -0.0039459 Median :4.013e+07 Median : 112.88 Median : 0.0005062 Median : 0.0004148 Mean :5.449e+08 Mean : 380.19 Mean : 0.0003592 Mean : 0.0003885 3rd Qu.:3.181e+08 3rd Qu.: 497.14 3rd Qu.: 0.0049885 3rd Qu.: 0.0046277 Max. :1.146e+10 Max. :1561.80 Max. : 0.0594595 Max. : 0.0540658 low_new close_new volume_new Adj.close_new Min. : -0.0821116 Min. : -0.0680141 Min. : -0.754927 Min. : -0.0680141 </pre>			

low_new	close_new	volume_new	Adj.close_new
Max. : 1.146e+10	Max. : 1561.80	Max. : 0.0594595	Max. : 0.0540658
Min. : -0.0821116	Min. : -0.0680141	Min. : -0.754927	Min. : -0.0680141
1st Qu.: -0.0041704	1st Qu.: -0.0044001	1st Qu.: -0.092642	1st Qu.: -0.0044001
Median : 0.0005606	Median : 0.0004455	Median : 0.004051	Median : 0.0004455
Mean : 0.0004167	Mean : 0.0004045	Mean : 0.017172	Mean : 0.0004045
3rd Qu.: 0.0047436	3rd Qu.: 0.0050338	3rd Qu.: 0.109569	3rd Qu.: 0.0050338
Max. : 0.1067194	Max. : 0.1078900	Max. : 2.996867	Max. : 0.1078900

b) For your samples, use the functions `mean()`, `max()`, `var()` and `quantile(, .25)` to compute the mean, maximum, variance and 1st quartile respectively for each column which has successive differences. Show your R code and the resulting values.

Do the same thing by using Excel. Show your Excel commands.

```

Console Terminal Jobs
F:/Data Science/DataScience_2019501111/Data Mining/Final exam/
> mean(data_1000$open_new)
[1] 0.0005955025
> mean(data_1000$high_new)
[1] 0.0004184797
> mean(data_1000$low_new)
[1] 0.0005022487
> mean(data_1000$close_new)
[1] 0.0003369592
> mean(data_1000$volume_new)
[1] 0.007551912
> mean(data_1000$Adj.close_new)
[1] 0.0003369592
> var(data_1000$open_new)
[1] 8.714339e-05
> var(data_1000$high_new)
[1] 6.119132e-05
> var(data_1000$low_new)
[1] 8.313995e-05
> var(data_1000$close_new)
[1] 7.637739e-05
> var(data_1000$volume_new)
[1] 0.0327711
> var(data_1000$Adj.close_new)
[1] 7.637739e-05
> max(data_1000$open_new)
[1] 0.1067121
> max(data_1000$high_new)
[1] 0.03439077
> max(data_1000$low_new)
[1] 0.09108332
> max(data_1000$close_new)
[1] 0.05732732
> max(data_1000$volume_new)
[1] 1.677175
> max(data_1000$Adj.close_new)
[1] 0.05732732
> quantile(data_1000$open_new, 0.25)
25%
-0.003961827
> quantile(data_1000$high_new, 0.25)
25%

```

```
Console Terminal x Jobs x
F:/Data Science/DataScience_2019501111/Data Mining/Final exam/

-0.003961827
> quantile(data_1000$high_new,0.25)
25%
-0.003443228
> quantile(data_1000$low_new,0.25)
25%
-0.003897353
> quantile(data_1000$close_new,0.25)
25%
-0.004251294
> quantile(data_1000$volume_new,0.25)
25%
-0.1056329
> quantile(data_1000$Adj.close_new,0.25)
25%
-0.004251294
>
>
> mean(data_3000$open_new)
[1] 0.0003591911
> mean(data_3000$high_new)
[1] 0.0003884621
> mean(data_3000$low_new)
[1] 0.0004167
> mean(data_3000$close_new)
[1] 0.0004044752
> mean(data_3000$volume_new)
[1] 0.0171718
> mean(data_3000$Adj.close_new)
[1] 0.0004044752
> var(data_3000$open_new)
[1] 8.509529e-05
> var(data_3000$high_new)
[1] 6.81047e-05
> var(data_3000$low_new)
[1] 8.768766e-05
> var(data_3000$close_new)
[1] 8.588174e-05
> var(data_3000$volume_new)
[1] 0.03939109
```

```
Console Terminal x Jobs x
F:/Data Science/DataScience_2019501111/Data Mining/Final exam/
[1] 6.81047e-05
> var(data_3000$low_new)
[1] 8.768766e-05
> var(data_3000$close_new)
[1] 8.588174e-05
> var(data_3000$volume_new)
[1] 0.03939109
> var(data_3000$Adj.close_new)
[1] 8.588174e-05
> max(data_3000$open_new)
[1] 0.05945946
> max(data_3000$high_new)
[1] 0.05406578
> max(data_3000$low_new)
[1] 0.1067194
> max(data_3000$close_new)
[1] 0.10789
> max(data_3000$volume_new)
[1] 2.996867
> max(data_3000$Adj.close_new)
[1] 0.10789
> quantile(data_3000$open_new,0.25)
25%
-0.003965834
> quantile(data_3000$high_new,0.25)
25%
-0.003945885
> quantile(data_3000$low_new,0.25)
25%
-0.004170403
> quantile(data_3000$close_new,0.25)
25%
-0.00440009
> quantile(data_3000$volume_new,0.25)
25%
-0.09264194
> quantile(data_3000$Adj.close_new,0.25)
25%
-0.00440009
>
```

c) Compute the same quantities in part b on the entire data set and show your answers. How much do they differ from your answers in part b? Do you find any significant difference between two sample values like mean in comparison with entire data? If so what explanation you can give for that?


Do the same thing by using Excel. Show your Excel commands.

```
Console Terminal x Jobs x
F:/Data Science/DataScience_2019501111/Data Mining/Final exam/

> mean(data$open_new)
[1] 0.000329528
> mean(data$high_new)
[1] 0.0003188991
> mean(data$low_new)
[1] 0.0003266191
> mean(data$close_new)
[1] 0.0003303709
> mean(data$volume_new)
[1] 0.02062874
> mean(data$Adj.close_new)
[1] 0.0003303709
> var(data$open_new)
[1] 9.027493e-05
> var(data$high_new)
[1] 6.939914e-05
> var(data$low_new)
[1] 8.646474e-05
> var(data$close_new)
[1] 9.350347e-05
> var(data$volume_new)
[1] 0.09080738
> var(data$Adj.close_new)
[1] 9.350347e-05
> max(data$open_new)
[1] 0.1067121
> max(data$high_new)
[1] 0.08037943
> max(data$low_new)
[1] 0.1067194
> max(data$close_new)
[1] 0.1158004
> max(data$volume_new)
[1] 26.51968
> max(data$Adj.close_new)
[1] 0.1158004
> quantile(data$open_new,0.25)
25%
-0.004110794
> quantile(data$high_new,0.25)
25%
-0.004110794
> quantile(data$open_new,0.25)
25%
-0.004110794
> quantile(data$high_new,0.25)
25%
-0.003772912
> quantile(data$low_new,0.25)
25%
-0.003996406
> quantile(data$close_new,0.25)
25%
-0.004121264
> quantile(data$volume_new,0.25)
25%
-0.09553922
> quantile(data$Adj.close_new,0.25)
25%
-0.004121264
> |
```

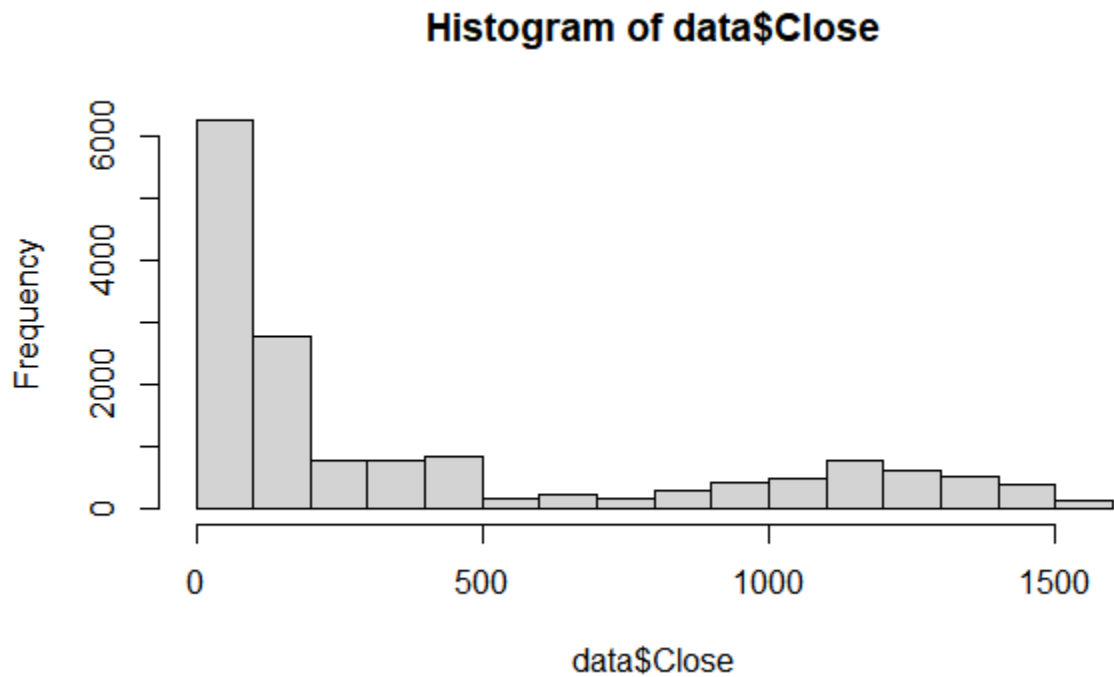
d) Use R to produce a single graph displaying a boxplot for open, close, high and low. Include the R commands and the plot.

Do the same thing by using Excel. Show your Excel commands



e) Use R to produce a frequency histogram for Close values. Use intervals of width 2000 beginning at 0. Include the R commands and the plot.

Do the same thing by using Excel. Show your Excel commands. (10+10=20M)



2. Implement Apriori Algorithm or use built in packages to find out the frequent itemsets and generate rules for frequent itemsets. Trace and submit the program output for the following given dataset of transactions with a minimum support of 3. (10M)

TID, Items
101, A,B,C,D,E
102, A,C,D
103, D,E
104, B,C,E
105, A,B,D,E
106, A,B
107, B,D,E
108, A,B,D
109, A,D
110, D,E


```
Console Terminal x Jobs x
F:/Data Science/DataScience_2019501111/Data Mining/Final exam/
> write.csv(data, "ItemList.csv", quote = FALSE, row.names = TRUE)
> transactions = read.transactions("ItemList.csv", sep=',', rm.duplicates = TRUE)
> inspect(transactions)
  items
[1] {Items}
[2] {1,A,B,C,D,E}
[3] {2,A,C,D}
[4] {3,D,E}
[5] {4,B,C,E}
[6] {5,A,B,D,E}
[7] {6,A,B}
[8] {7,B,D,E}
[9] {8,A,B,D}
[10] {9,A,D}
[11] {10,D,E}
>
> freqItemsets <- apriori(transactions, parameter = list(sup = 0.03, conf = 0.5, target
="frequent itemsets"))
Apriori

Parameter specification:
confidence minval smax arem aval originalsupport maxtime support minlen maxlen
      NA      0.1      1 none FALSE              TRUE        5      0.03      1      10
      target ext
frequent itemsets TRUE

Algorithmic control:
filter tree heap memopt load sort verbose
  0.1 TRUE TRUE  FALSE TRUE    2    TRUE

Absolute minimum support count: 0

set item appearances ...[0 item(s)] done [0.00s].
set transactions ...[16 item(s), 11 transaction(s)] done [0.00s].
sorting and recoding items ... [16 item(s)] done [0.00s].
creating transaction tree ... done [0.00s].
checking subsets of size 1 2 3 4 5 6 done [0.00s].
sorting transactions ... done [0.00s].
writing ... [128 set(s)] done [0.00s].
creating s4 object ... done [0.06s].
```

3. Build Decision Trees by using i) information gain and ii) misclassification error rate for Lenses Data Set provided at <http://archive.ics.uci.edu/ml/datasets/Lenses>. In terms of tree size what do you conclude comparing these two? (10M)

```
Console Terminal x Jobs x
F:/Data Science/DataScience_2019501111/Data Mining/Final exam/

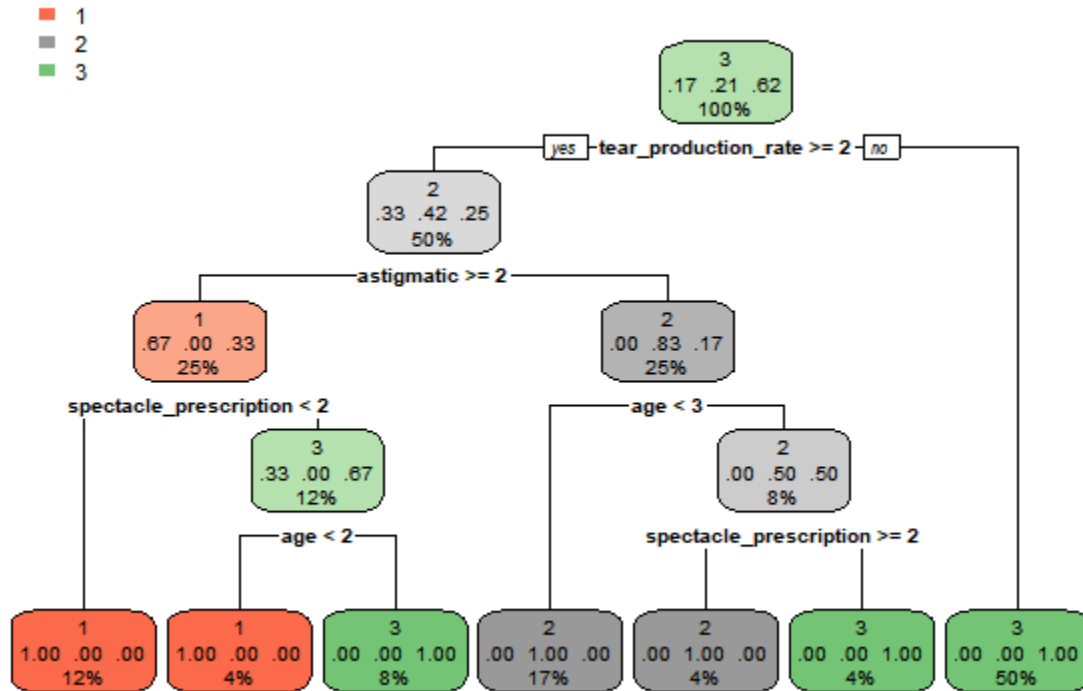
The downloaded binary packages are in
  C:\Users\hp\AppData\Local\Temp\Rtmp0qk3yh\downloaded_packages
> install.packages("rpart.plot")
WARNING: Rtools is required to build R packages but is not currently installed. Please
  download and install the appropriate version of Rtools before proceeding:

https://cran.rstudio.com/bin/windows/Rtools/
Installing package into 'C:/Users/hp/Documents/R/win-library/4.0'
(as 'lib' is unspecified)
trying URL 'https://cran.rstudio.com/bin/windows/contrib/4.0/rpart.plot_3.0.9.zip'
Content type 'application/zip' length 1033909 bytes (1009 KB)
downloaded 1009 KB

package 'rpart.plot' successfully unpacked and MD5 sums checked

The downloaded binary packages are in
  C:\Users\hp\AppData\Local\Temp\Rtmp0qk3yh\downloaded_packages
> setwd("F:\\Data Science\\DataScience_2019501111\\Data Mining\\Final exam")
> lensdata = read.csv("lenses.data.csv", header = FALSE, col.names = c("index", "age",
  "spectacle_prescription", "astigmatic", "tear_production_rate", "class"))
>
> lensdata$index <- NULL
>
> library(rpart)
warning message:
package 'rpart' was built under R version 4.0.3
> y<-as.factor(lensdata[,5])
> x<-lensdata[,1:4]
>
> model1<-rpart(y~.,x, parms = list(split = 'information'),
  + control=rpart.control(minsplit=0,minbucket=0,cp=-1, maxcompete=0, maxsu
  rrogate=0, usesurrogate=0, xval=0,maxdepth=5))
>
>
> library(rpart.plot)
warning message:
package 'rpart.plot' was built under R version 4.0.3
> rpart.plot(model1)
> |
```

```
>
>
> library(rpart.plot)
> rpart.plot(model1)
>
> #Information gain
> gain <- sum(y==predict(model1,x,type="class"))/length(y)
> gain
[1] 1
>
> #misclassification
> error_rate <- 1-sum(y==predict(model1,x,type="class"))/length(y)
> error_rate
[1] 0
> |
```



4. Fit 1, 2 and 3-nearest-neighbor classifiers to the Liver Disorders Data Set at <http://archive.ics.uci.edu/ml/datasets/Liver+Disorders> for measures Euclidean and cosine. Last but one column is a decision attribute. Replace decision values in to 4 classes ($0 \leq c_1 < 5$, $5 \leq c_2 < 10$, $10 \leq c_3 < 15$, $15 \leq c_4 \leq 20$). Last column is a data split column in to training and test sets. 1 means the object is used for training. 2 means the object is used for testing. Explain the input parameters you provided for the classifier. Compute the misclassification error on the training data and also on the test data. Annotate your program. (10M)

```
Source
Console Terminal x Jobs x
F:/Data Science/DataScience_2019501111/Data Mining/Final exam/
> setwd("F:\\Data Science\\DataScience_2019501111\\Data Mining\\Final exam")
> data = read.csv("Liver_Data.csv", header = FALSE, col.names = c("mcv", "alkphos", "sg
pt", "sgot", "gammagt", "drinks", "selector"))
>
> #converting the decision attribute into classes
> data$drinks = cut(data$drinks, breaks = c(0,5,10,15,20,25), labels = c('c1', 'c2', 'c
3', 'c4', 'c4'), right = FALSE)
> data = na.omit(data)
>
> #traing and test sets
> traindata = subset(data, data$selector == 1)
> testdata = subset(data, data$selector == 2)
>
> x_train <- subset(traindata, select = -c(selector, drinks))
> x_test <- subset(testdata, select = -c(selector, drinks))
>
> y_train = traindata[,6, drop = TRUE]
> y_test = testdata[,6, drop = TRUE]
>
>
> #For Training Data
> #knn for k=1
> library(class)
> model1 = knn(x_train, x_test, y_train, k = 1)
> 1-sum(y_train==model1)/length(y_train) # 0
[1] 0.2827586
warning messages:
1: In `==.default`(y_train, model1) :
  longer object length is not a multiple of shorter object length
2: In is.na(e1) | is.na(e2) :
  longer object length is not a multiple of shorter object length
>
> #knn for k=2
> model2 = knn(x_train, x_train, y_train, k = 2)
> 1-sum(y_train==model2)/length(y_train) # 0.1586207
[1] 0.1655172
>
> #knn for k=3
> model3 = knn(x_train, x_train, y_train, k = 3)
> 1-sum(y_train==model3)/length(y_train) # 0.2137931
[1] 0.2068866
```

```

>
> #knn for k=2
> model2 = knn(x_train, x_train, y_train, k = 2)
> 1-sum(y_train==model2)/length(y_train) # 0.1586207
[1] 0.1655172
>
> #knn for k=3
> model3 = knn(x_train, x_train, y_train, k = 3)
> 1-sum(y_train==model3)/length(y_train) # 0.2137931
[1] 0.2068966
>
>
> #For Test Data
> #knn for k=1
> model4 = knn(x_train, x_test, y_train, k = 1)
> 1-sum(y_test==model4)/length(y_test) # 0.44
[1] 0.445
>
> #knn for k=2
> model5 = knn(x_train, x_test, y_train, k = 2)
> 1-sum(y_test==model5)/length(y_test) # 0.42
[1] 0.445
>
> #knn for k=3
> model6 = knn(x_train, x_test, y_train, k = 3)
> 1-sum(y_test==model6)/length(y_test) # 0.405
[1] 0.4
> |

```

5. Use Support Vector machine for above problem. And compare the performance of both. Explain the input parameters you provided for the classifier. (10M)

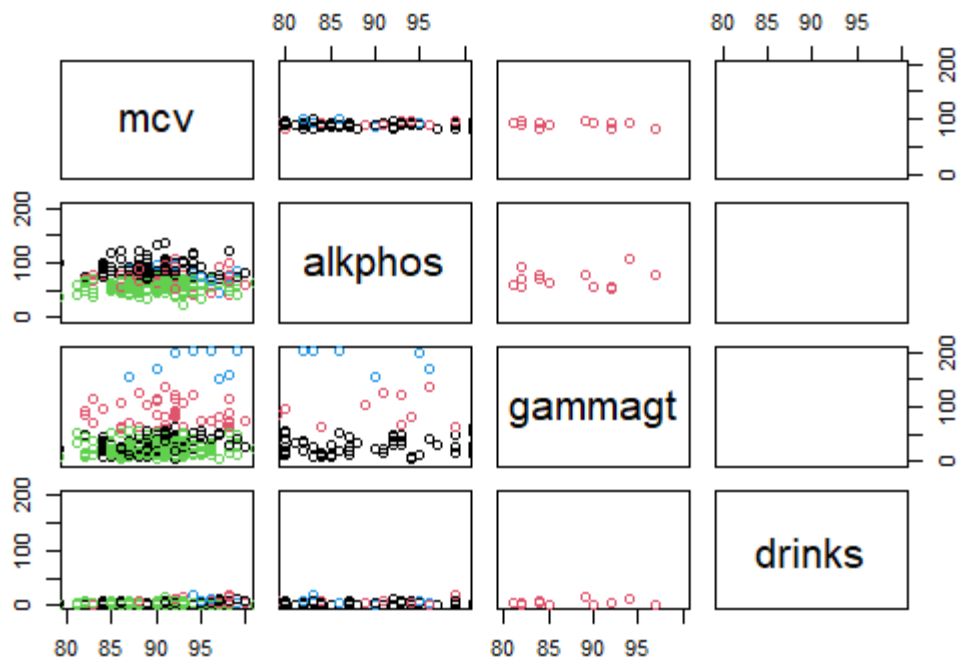
```

package 'e1071' successfully unpacked and MD5 sums checked

The downloaded binary packages are in
  C:\Users\hp\AppData\Local\Temp\Rtmp0qk3yh\downloaded_packages
> setwd("F:\\Data Science\\DataScience_2019501111\\Data Mining\\Final exam")
> data = read.csv("Liver_data.csv", header = FALSE, col.names = c("mcv", "alkphos", "sg
pt", "sgot", "gammagt", "drinks", "selector"))
>
> #converting the decision attribute into classes
> data$drinks = cut(data$drinks, breaks = c(0,5,10,15,20,25), labels = c('c1', 'c2', 'c
3', 'c4', 'c4'), right = FALSE)
> data = na.omit(data)
>
> #traing and test sets
> traindata = subset(data, data$selector == 1)
> testdata = subset(data, data$selector == 2)
>
> x_train <- subset(traindata, select = -c(selector, drinks))
> x_test <- subset(testdata, select = -c(selector, drinks))
>
> y_train = traindata[,6, drop = TRUE]
> y_test = testdata[,6, drop = TRUE]
>
> library(e1071)
warning message:
package 'e1071' was built under R version 4.0.3
>
> #For training
> model = svm(x_train, y_train)
> 1-sum(y_train==predict(model,x_train))/length(y_train) # 0.2137931
[1] 0.2137931
>
> #For test data
> 1-sum(y_test==predict(model,x_test))/length(y_test) # 0.285
[1] 0.285
>
> #The misclassification error is high for KNN so, we can prefer SVM over KNN
> |

```

6. Create k-means clusters for k=4 for the Liver Disorders Data Set at <http://archive.ics.uci.edu/ml/datasets/Liver+Disorders> . Explain the input parameters you provided for the clustering algorithm. Plot the fitted cluster centers using a different color. Finally assign the cluster membership for the points to the nearest cluster center. Color the points according to their cluster membership. (10+10=20M)



7. Compute the misclassification error that would result if you used your clustering rule to classify the data by assigning the majority class of the cluster. (10M)

8. Consider the dataset BSE_Sensex_Index.csv. Create an extra column of successive growth rate for column close where the successive growth rate is defined as $(\text{value of day } x - \text{value of day } x-1) / \text{value of day } x-1$. Use a z score cut off of 3 to identify any outliers. List the respective dates from the csv file on which day these outliers fall. (10M)

```

Source
Console Terminal x Jobs x
F:/Data Science/DataScience_2019501111/Data Mining/Final exam/
> data = read.csv("BSE_Sensex_Index.csv", header = TRUE)
>
> view(data)
> summary(data)
      Date      open      high      low
Length:15447   Min.   : 16.66   Min.   : 16.66   Min.   : 16.66
Class :character 1st Qu.: 79.98   1st Qu.: 80.72   1st Qu.: 79.39
Mode  :character Median : 115.97   Median : 117.01   Median : 114.85
      Mean : 393.96   Mean : 396.59   Mean : 391.19
      3rd Qu.: 619.74   3rd Qu.: 621.40   3rd Qu.: 616.46
      Max.   :1564.98   Max.   :1576.09   Max.   :1555.46
      close      volume      Adj.Close
Min.   : 16.66   Min.   :6.800e+05   Min.   : 16.66
1st Qu.: 79.98   1st Qu.:5.830e+06   1st Qu.: 79.98
Median : 116.00   Median :4.326e+07   Median : 116.00
Mean : 394.05   Mean :5.864e+08   Mean : 394.05
3rd Qu.: 620.07   3rd Qu.:3.832e+08   3rd Qu.: 620.07
Max.   :1565.15   Max.   :1.146e+10   Max.   :1565.15
> data$Date = as.Date(data$Date, format='%m/%d/%Y')
>
> successive_difference <- function(x) {
+   n = length(x)
+   for (i in 1:(length(x))) {
+     x[i] <- (x[i] - x[i + 1]) / x[i + 1]
+   }
+   x[length(x)] = (x[length(x) - 1] + x[length(x) - 2] + x[length(x) - 3]) / 3
+   return(x)
+ }
>
> data$successive_growth <- successive_difference(data$close)
>
>
> #calculating z-scores
> sgrmean <- mean(data$successive_growth, na.rm=TRUE)
> sgrsd <- sd(data$successive_growth, na.rm=TRUE)
> z<-(data$successive_growth - sgrmean) / sgrsd
> sort(z)
[1] -21.200164 -9.377746 -9.268692 -9.141750 -8.595889 -7.911031 -7.134352
[8] -7.067886 -7.033645 -6.975723 -6.937815 -6.878679 -6.360330 -6.358610
[15] -6.343812 -6.061018 -5.969685 -5.599374 -5.496172 -5.481801 -5.400803
[22] -5.378888 -5.333345 -5.133836 -5.114868 -5.086856 -4.888354 -4.888753

```



```
Source
Console Terminal x Jobs x
F:/Data Science/DataScience_2019501111/Data Mining/Final exam/
[785] -1.504077 -1.503930 -1.503767 -1.502479 -1.502254 -1.501164 -1.499653
[792] -1.498687 -1.496642 -1.496166 -1.495191 -1.493551 -1.493444 -1.493292
[799] -1.491062 -1.490834 -1.490794 -1.489958 -1.485868 -1.484670 -1.483748
[806] -1.483516 -1.482876 -1.482448 -1.482356 -1.481304 -1.480835 -1.480189
[813] -1.480096 -1.479982 -1.479827 -1.479415 -1.478708 -1.478207 -1.477655
[820] -1.477494 -1.477464 -1.475556 -1.474689 -1.473455 -1.472859 -1.472644
[827] -1.472447 -1.472386 -1.471786 -1.471159 -1.470927 -1.470546 -1.470337
[834] -1.470054 -1.468870 -1.468867 -1.467686 -1.466569 -1.466097 -1.465232
[841] -1.464190 -1.462344 -1.461426 -1.461369 -1.461190 -1.460763 -1.460113
[848] -1.457398 -1.457349 -1.456155 -1.455739 -1.455686 -1.455235 -1.455198
[855] -1.455016 -1.454341 -1.453796 -1.453590 -1.453399 -1.453068 -1.452906
[862] -1.451952 -1.450818 -1.449287 -1.447390 -1.446905 -1.446385 -1.445375
[869] -1.445363 -1.445034 -1.445027 -1.444987 -1.441605 -1.440731 -1.440410
[876] -1.440146 -1.439633 -1.436170 -1.436094 -1.435973 -1.433590 -1.433567
[883] -1.433210 -1.433186 -1.432202 -1.431674 -1.430946 -1.428230 -1.427024
[890] -1.426090 -1.425999 -1.424783 -1.423981 -1.423193 -1.422761 -1.422607
[897] -1.422158 -1.422094 -1.421540 -1.420730 -1.420679 -1.419969 -1.419312
[904] -1.417832 -1.416308 -1.416218 -1.415207 -1.414047 -1.413762 -1.413268
[911] -1.411751 -1.410838 -1.410105 -1.408879 -1.408103 -1.407752 -1.406047
[918] -1.404931 -1.403099 -1.402938 -1.400542 -1.400333 -1.399500 -1.399271
[925] -1.398907 -1.398231 -1.396046 -1.394765 -1.394478 -1.393779 -1.393586
[932] -1.393435 -1.393284 -1.393194 -1.392305 -1.389947 -1.389915 -1.389368
[939] -1.388574 -1.388493 -1.388149 -1.388138 -1.386722 -1.386128 -1.386004
[946] -1.385513 -1.385181 -1.385039 -1.383151 -1.382864 -1.382830 -1.382830
[953] -1.382748 -1.379377 -1.379336 -1.379104 -1.378806 -1.378641 -1.376414
[960] -1.375134 -1.374976 -1.374949 -1.374707 -1.374220 -1.374176 -1.373894
[967] -1.372775 -1.371636 -1.371541 -1.371007 -1.369255 -1.369064 -1.367802
[974] -1.366266 -1.365630 -1.364852 -1.362457 -1.361316 -1.361288 -1.360593
[981] -1.359933 -1.358861 -1.358227 -1.358002 -1.357692 -1.356240 -1.355144
[988] -1.353995 -1.352105 -1.351562 -1.349644 -1.349380 -1.349251 -1.348937
[995] -1.348035 -1.347070 -1.346192 -1.346087 -1.346066 -1.345068
[ reached getoption("max.print") -- omitted 14447 entries ]
> data$zscores <- z
>
> #Dates of the outliers
> dates<-subset(data[,1],data[, "zscores"] >= 3.0 | data[, "zscores"] <= -3.0)
> view(dates)
>
> write.csv(dates, "OutliersDatesData_.csv", quote = FALSE, row.names = TRUE)
> |
```